

Wild Blueberry Yield Prediction

Upload Daataset file

from google.colab import files
files.upload()

Choose files

WildBlueber...tionData.csv

- WildBlueberryPollinationSimulationData.csv(text/csv) - 85259 bytes, last modified: 15/02/2024 - 100% done

Saving WildBlueberryPollinationSimulationData.csv to WildBlueberryPollinationSimulationData.csv

```
{'WildBlueberryPollinationSimulationData.csv':  
b'Row#,clonesize,honeybee,bumbles,andrena,osmia,MaxOfUpperTRange,MinOfUpperTRange,AverageOfUpperTRange,MaxOfLowerTRange,MinOfLowerTRange,AverageOfLowerTRange,RainingDays,AverageRainingDays,fruitset'
```

import pandas as pd

Provide the path to your dataset file
file_path = "/content/WildBlueberryPollinationSimulationData.csv"

Read the dataset into a Pandas DataFrame
df = pd.read_csv(file_path)

Display the first few rows of the DataFrame
df.head()

	Row#	clonesize	honeybee	bumbles	andrena	osmia	MaxOfUpperTRange	MinOfUpperTRange	AverageOfUpperTRange	MaxOfLowerTRange	MinOfLowerTRange	AverageOfLowerTRange	RainingDays	AverageRainingDays	fruitset
0	0	37.5	0.75	0.25	0.25	0.25	86.0	52.0	71.9	62.0	30.0	50.8	16.0		(
1	1	37.5	0.75	0.25	0.25	0.25	86.0	52.0	71.9	62.0	30.0	50.8	1.0		(
2	2	37.5	0.75	0.25	0.25	0.25	94.6	57.2	79.0	68.2	33.0	55.9	16.0		(
3	3	37.5	0.75	0.25	0.25	0.25	94.6	57.2	79.0	68.2	33.0	55.9	1.0		(
4	4	37.5	0.75	0.25	0.25	0.25	86.0	52.0	71.9	62.0	30.0	50.8	24.0		(

Next steps: [View recommended plots](#)

Task 1: Exploratory Data Analysis

print(df.columns)

```
Index(['Row#', 'clonesize', 'honeybee', 'bumbles', 'andrena', 'osmia',  
      'MaxOfUpperTRange', 'MinOfUpperTRange', 'AverageOfUpperTRange',  
      'MaxOfLowerTRange', 'MinOfLowerTRange', 'AverageOfLowerTRange',  
      'RainingDays', 'AverageRainingDays', 'fruitset', 'fruitmass', 'seeds',  
      'yield'],  
      dtype='object')
```

print(df['yield'].describe())

```
count      777.000000  
mean      6012.849165  
std       1356.955318  
min       1637.704022  
25%       5124.854901  
50%       6107.382466  
75%       7022.189731  
max       8969.401842  
Name: yield, dtype: float64
```

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt


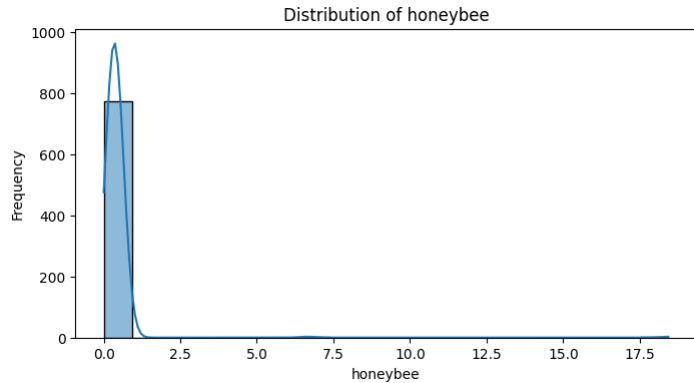
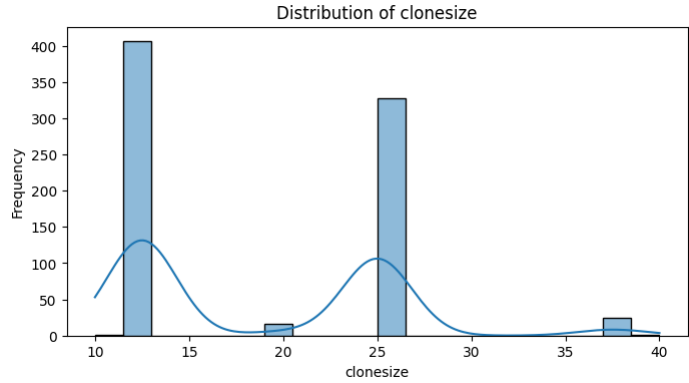
Display the first few rows of the DataFrame
print(df.head())


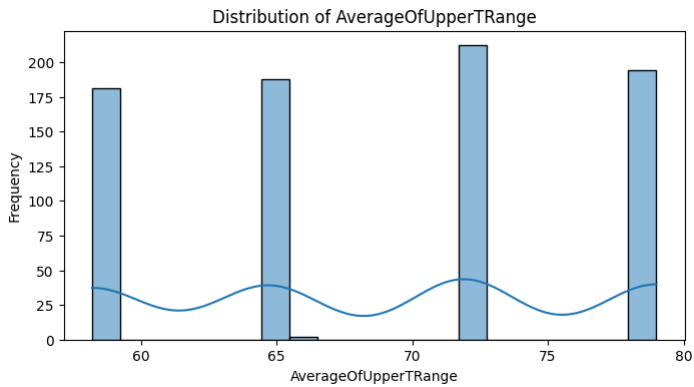
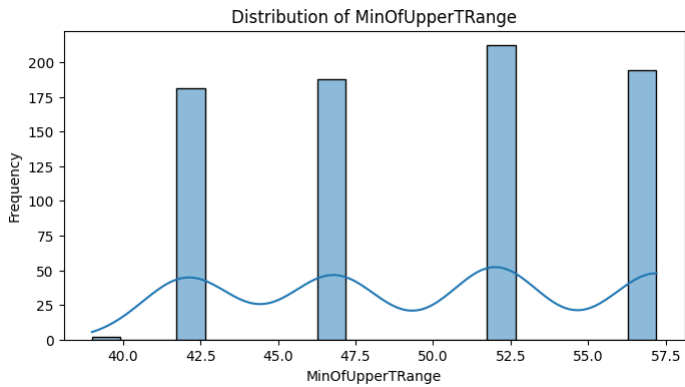
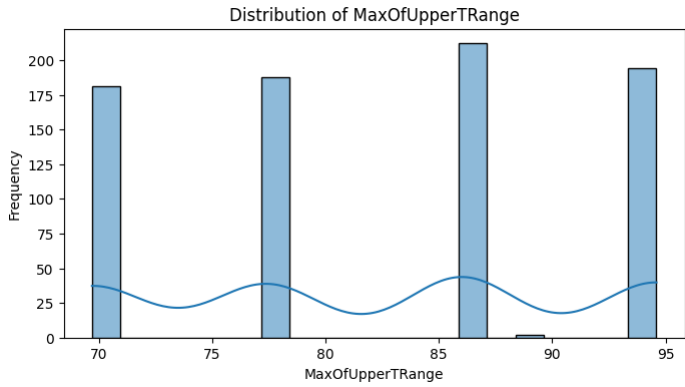
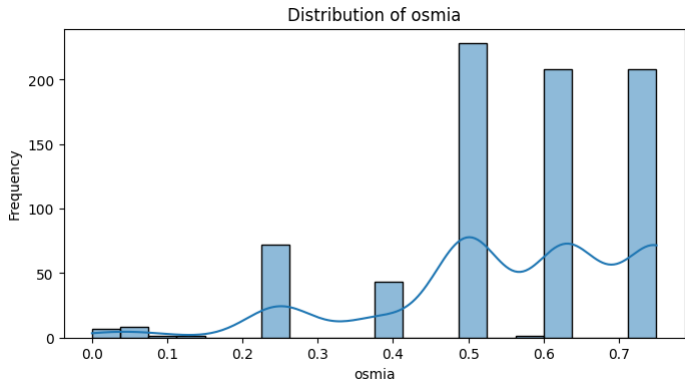
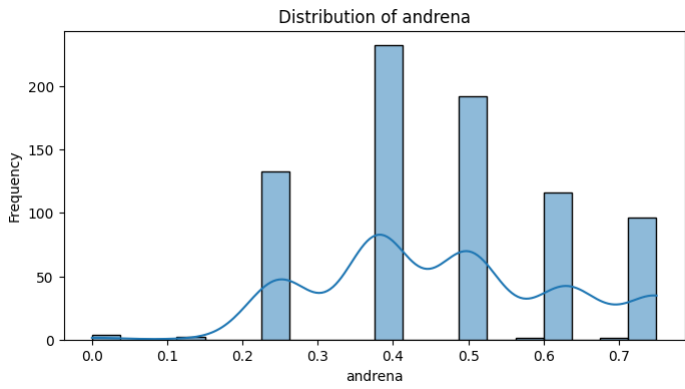

Display summary statistics
print(df.describe())


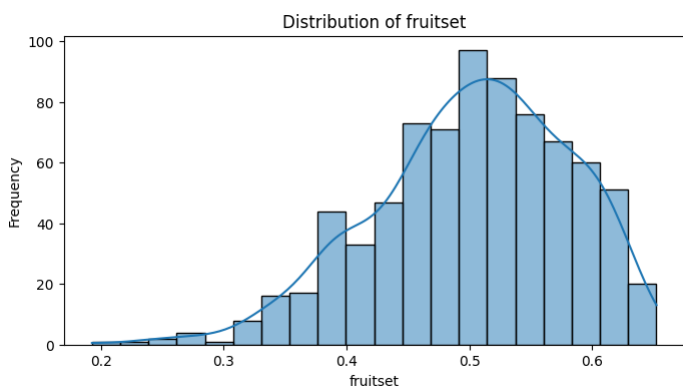
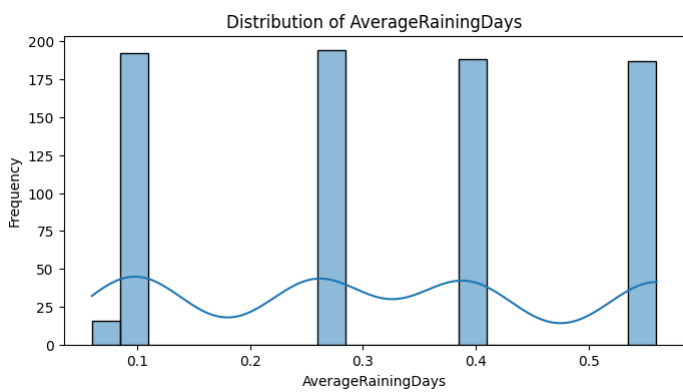
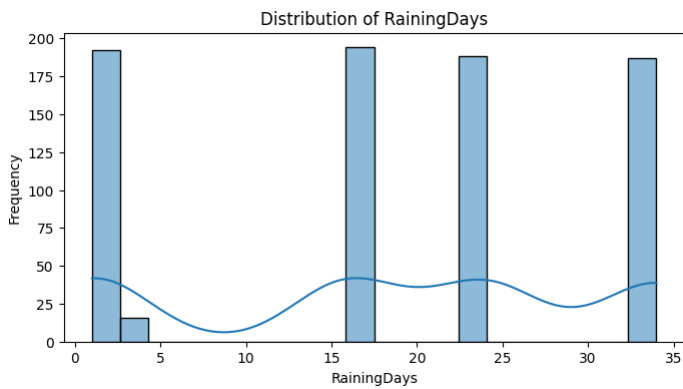
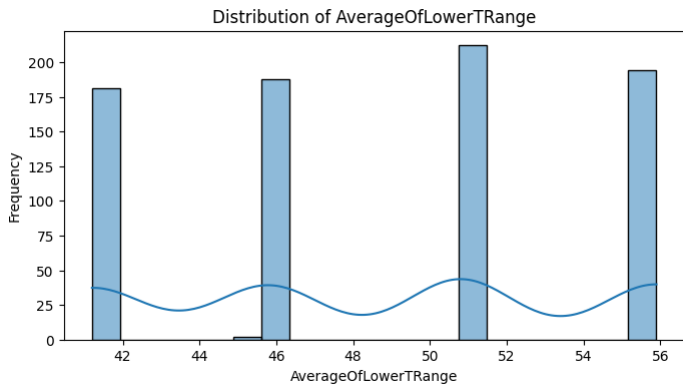
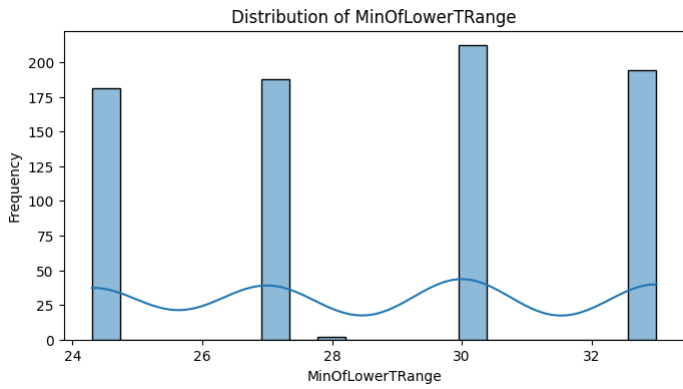

Visualize the distribution of each numerical variable




```
numerical_columns = df.select_dtypes(include=['float64']).columns  
for column in numerical_columns:  
    plt.figure(figsize=(8, 4))  
    sns.histplot(df[column], bins=20, kde=True)  
    plt.title(f'Distribution of {column}')  
    plt.xlabel(column)  
    plt.ylabel('Frequency')  
    plt.show()
```

Row#	clonesize	honeybee	bumbles	andrena	osmia	MaxOfUpperTRange	\
0	0	37.5	0.75	0.25	0.25	0.25	86.0
1	1	37.5	0.75	0.25	0.25	0.25	86.0
2	2	37.5	0.75	0.25	0.25	0.25	94.6
3	3	37.5	0.75	0.25	0.25	0.25	94.6
4	4	37.5	0.75	0.25	0.25	0.25	86.0
	MinOfUpperTRange	AverageOfUpperTRange	MaxOfLowerTRange	MinOfLowerTRange	\		
0	52.0		71.9	62.0	30.0		
1	52.0		71.9	62.0	30.0		
2	57.2		79.0	68.2	33.0		
3	57.2		79.0	68.2	33.0		
4	52.0		71.9	62.0	30.0		
	AverageOfLowerTRange	RainingDays	AverageRainingDays	fruitset	fruitmass	\	
0	50.8	16.0		0.26	0.410652	0.408159	
1	50.8	1.0		0.10	0.444254	0.425458	
2	55.9	16.0		0.26	0.383787	0.399172	
3	55.9	1.0		0.10	0.407564	0.408789	
4	50.8	24.0		0.39	0.354413	0.382703	
	seeds	yield					
0	31.678898	3813.165795					
1	33.449385	4947.605663					
2	30.546306	3866.798965					
3	31.562586	4303.943030					
4	28.873714	3436.493543					
Row#	clonesize	honeybee	bumbles	andrena	osmia	\	
count	777.000000	777.000000	777.000000	777.000000	777.000000	777.000000	
mean	388.000000	18.767696	0.417133	0.282389	0.468817	0.562062	
std	224.444871	6.999063	0.978904	0.066343	0.161052	0.169119	
min	0.000000	10.000000	0.000000	0.000000	0.000000	0.000000	
25%	194.000000	12.500000	0.250000	0.250000	0.380000	0.500000	
50%	388.000000	12.500000	0.250000	0.250000	0.500000	0.630000	
75%	582.000000	25.000000	0.500000	0.380000	0.630000	0.750000	
max	776.000000	40.000000	18.430000	0.585000	0.750000	0.750000	
	MaxOfUpperTRange	MinOfUpperTRange	AverageOfUpperTRange	\			
count	777.000000	777.000000	777.000000				
mean	82.277091	49.700515	68.723037				
std	9.193745	5.595769	7.676984				
min	69.700000	39.000000	58.200000				
25%	77.400000	46.800000	64.700000				
50%	86.000000	52.000000	71.900000				
75%	89.000000	52.000000	71.900000				
max	94.600000	57.200000	79.000000				
	MaxOfLowerTRange	MinOfLowerTRange	AverageOfLowerTRange	RainingDays	\		
count	777.000000	777.000000	777.000000	777.000000			
mean	59.309395	28.690219	48.613127	18.309292			
std	6.647760	3.209547	5.417072	12.124226			
min	50.200000	24.300000	41.200000	1.000000			
25%	55.800000	27.000000	45.800000	3.770000			
50%	62.000000	30.000000	50.800000	16.000000			
75%	66.000000	30.000000	50.800000	24.000000			
max	68.200000	33.000000	55.900000	34.000000			
	AverageRainingDays	fruitset	fruitmass	seeds	yield		
count	777.000000	777.000000	777.000000	777.000000	777.000000		
mean	0.320000	0.502121	0.445983	36.122432	6012.849165		
std	0.171279	0.079445	0.040333	4.377889	1356.955318		
min	0.060000	0.192732	0.311921	22.079199	1637.704022		
25%	0.100000	0.454725	0.416281	33.116091	5124.854901		
50%	0.260000	0.508297	0.445587	36.166044	6107.382466		
75%	0.390000	0.561297	0.476149	39.239668	7022.189731		
max	0.560000	0.652144	0.535660	46.585105	8969.401842		










```
# Correlation matrix
correlation_matrix = df.corr()

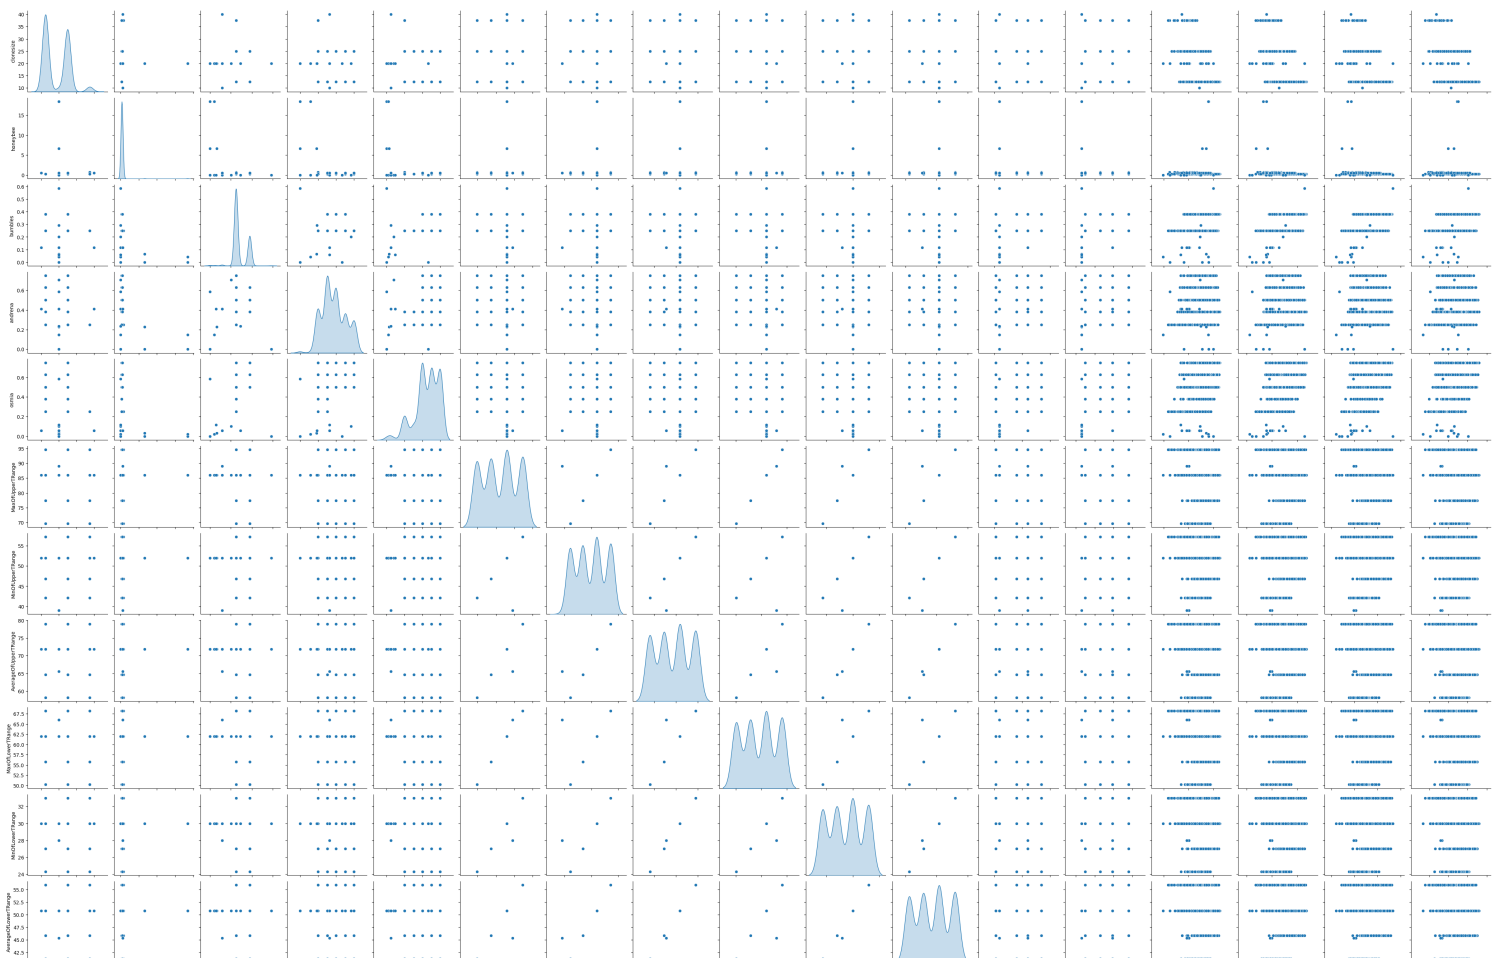
# Heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix")
plt.show()
```



With Heatmaps we can identify patterns of correlation or dependence between different variables in the DataFrame. Positive correlations are indicated by warmer colors, negative correlations by cooler colors, and the intensity of the color represents the strength of the correlation. The annotated values in each cell provide the exact correlation coefficient between the corresponding variables.


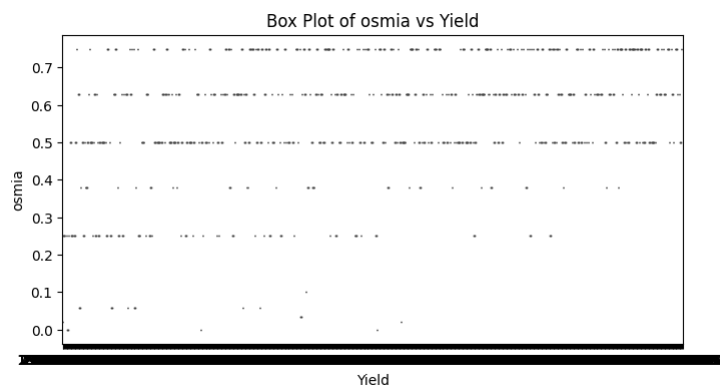

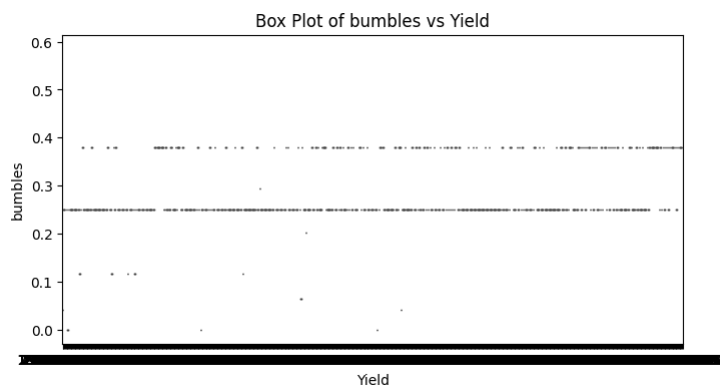


```
# Pair plots
sns.pairplot(df, vars=numerical_columns, diag_kind='kde')
plt.suptitle("Pair Plots of Numerical Variables", y=1.02)
plt.show()
```







Pair Plots of Numerical Variables









The resulting visualization is a matrix of scatterplots, where each cell shows the relationship between two numerical variables. The diagonal contains kernel density plots for each variable, and the scatterplots in the lower and upper triangles show the bivariate relationships. This type of visualization is useful for identifying patterns, trends, and potential correlations between different numerical features in the dataset.

```
# Box plots
for column in numerical_columns:
    plt.figure(figsize=(8, 4))
    sns.boxplot(x='yield', y=column, data=df)
    plt.title(f'Box Plot of {column} vs Yield')
    plt.xlabel('Yield')
    plt.ylabel(column)
    plt.show()
```







The resulting visualizations are a series of box plots, each showing the distribution of a numerical variable with respect to the 'yield'. Box plots provide information about the median, quartiles, and potential outliers in the distribution of each variable for different levels of the 'yield'.

Task 2: Preprocessing

```
# Check for missing values
missing_values = df.isnull().sum()

# Print missing values
print("Missing Values:")
print(missing_values)

# Handle missing values (example: fill with mean)
df.fillna(df.mean(), inplace=True)

# Print the updated dataframe
print("\nUpdated DataFrame after handling missing values:")
print(df.head())
```

Missing Values:

Row#	0
clonesize	0
honeybee	0
bumbles	0
andrena	0
osmia	0
MaxOfUpperTRange	0
MinOfUpperTRange	0
AverageOfUpperTRange	0
MaxOfLowerTRange	0
MinOfLowerTRange	0
AverageOfLowerTRange	0
RainingDays	0
AverageRainingDays	0
fruitset	0
fruitmass	0
seeds	0
yield	0

dtype: int64

Updated DataFrame after handling missing values:

Row#	clonesize	honeybee	bumbles	andrena	osmia	MaxOfUpperTRange	MinOfUpperTRange	AverageOfUpperTRange	MaxOfLowerTRange	MinOfLowerTRange	AverageOfLowerTRange	RainingDays	AverageRainingDays	fruitset	fruitmass	seeds	yield
0	0	37.5	0.75	0.25	0.25	86.0	52.0	57.2	62.0	30.0	50.8	16.0	0.26	0.410652	0.408159	31.678898	3813.165795
1	1	37.5	0.75	0.25	0.25	86.0	52.0	57.2	62.0	30.0	50.8	1.0	0.10	0.444254	0.425458	33.449385	4947.605663
2	2	37.5	0.75	0.25	0.25	94.6	57.2	55.9	68.2	33.0	55.9	16.0	0.26	0.383787	0.399172	30.546306	3866.798965
3	3	37.5	0.75	0.25	0.25	94.6	57.2	55.9	68.2	33.0	55.9	1.0	0.10	0.407564	0.408789	31.562586	4303.943030
4	4	37.5	0.75	0.25	0.25	86.0	52.0	57.2	62.0	30.0	50.8	24.0	0.39	0.354413	0.382703	28.873714	3436.493543

```
print(df.columns)
```

```
Index(['Row#', 'clonesize', 'honeybee', 'bumbles', 'andrena', 'osmia',
      'MaxOfUpperTRange', 'MinOfUpperTRange', 'AverageOfUpperTRange',
      'MaxOfLowerTRange', 'MinOfLowerTRange', 'AverageOfLowerTRange',
      'RainingDays', 'AverageRainingDays', 'fruitset', 'fruitmass', 'seeds',
      'yield'],
      dtype='object')

from sklearn.preprocessing import LabelEncoder

# List of columns to encode
columns_to_encode = ['Row#', 'clonesize', 'honeybee', 'bumbles', 'andrena', 'osmia',
                    'MaxOfUpperTRange', 'MinOfUpperTRange', 'AverageOfUpperTRange',
                    'MaxOfLowerTRange', 'MinOfLowerTRange', 'AverageOfLowerTRange',
                    'RainingDays', 'AverageRainingDays', 'fruitset', 'fruitmass', 'seeds', 'yield']

# Instantiate LabelEncoder
label_encoder = LabelEncoder()

# Apply label encoding to each column
for column in columns_to_encode:
    df[column] = label_encoder.fit_transform(df[column])

# Apply one-hot encoding to each column
df = pd.get_dummies(df, columns=columns_to_encode)

# Display the first few rows of the DataFrame after one-hot encoding
print(df.head())

# Check the data types of the columns
print(df.dtypes)

# Check unique values for each column
for column in df.columns:
    print(f"Unique values in {column}: {df[column].unique()}")
```

```
Unique values in yield_726: [0 1]
Unique values in yield_727: [0 1]
Unique values in yield_728: [0 1]
Unique values in yield_729: [0 1]
Unique values in yield_730: [0 1]
Unique values in yield_731: [0 1]
Unique values in yield_732: [0 1]
Unique values in yield_733: [0 1]
Unique values in yield_734: [0 1]
Unique values in yield_735: [0 1]
Unique values in yield_736: [0 1]
Unique values in yield_737: [0 1]
Unique values in yield_738: [0 1]
Unique values in yield_739: [0 1]
Unique values in yield_740: [0 1]
Unique values in yield_741: [0 1]
Unique values in yield_742: [0 1]
Unique values in yield_743: [0 1]
Unique values in yield_744: [0 1]
Unique values in yield_745: [0 1]
Unique values in yield_746: [0 1]
Unique values in yield_747: [0 1]
Unique values in yield_748: [0 1]
Unique values in yield_749: [0 1]
Unique values in yield_750: [0 1]
Unique values in yield_751: [0 1]
Unique values in yield_752: [0 1]
Unique values in yield_753: [0 1]
Unique values in yield_754: [0 1]
Unique values in yield_755: [0 1]
Unique values in yield_756: [0 1]
Unique values in yield_757: [0 1]
Unique values in yield_758: [0 1]
Unique values in yield_759: [0 1]
Unique values in yield_760: [0 1]
Unique values in yield_761: [0 1]
Unique values in yield_762: [0 1]
Unique values in yield_763: [0 1]
Unique values in yield_764: [0 1]
Unique values in yield_765: [0 1]
Unique values in yield_766: [0 1]
Unique values in yield_767: [0 1]
Unique values in yield_768: [0 1]
Unique values in yield_769: [0 1]
Unique values in yield_770: [0 1]
Unique values in yield_771: [0 1]
Unique values in yield_772: [0 1]
Unique values in yield_773: [0 1]
Unique values in yield_774: [0 1]
Unique values in yield_775: [0 1]
Unique values in yield_776: [0 1]
```

```
# Check the columns in the DataFrame
print(df.columns)

Index(['Row#_0', 'Row#_1', 'Row#_2', 'Row#_3', 'Row#_4', 'Row#_5', 'Row#_6',
      'Row#_7', 'Row#_8', 'Row#_9',
      ...,
      'yield_767', 'yield_768', 'yield_769', 'yield_770', 'yield_771',
      'yield_772', 'yield_773', 'yield_774', 'yield_775', 'yield_776'],
      dtype='object', length=3972)
```

```
# Separate features and target variables
X = df.drop(columns=df.filter(like='yield').columns) # Exclude the one-hot encoded 'yield' columns
y = df.filter(like='yield') # Include all the new one-hot encoded 'yield' columns
```

```
# Print the first few rows of features and target variables
print("Features:")
print(X.head())
```

```
print("\nTarget:")
print(y.head())
```

```
Features:
Row#_0  Row#_1  Row#_2  Row#_3  Row#_4  Row#_5  Row#_6  Row#_7  Row#_8  \
0      1      0      0      0      0      0      0      0      0
1      0      1      0      0      0      0      0      0      0
2      0      0      1      0      0      0      0      0      0
3      0      0      0      1      0      0      0      0      0
4      0      0      0      0      1      0      0      0      0
```

```
Row#_9  ...  seeds_767  seeds_768  seeds_769  seeds_770  seeds_771  \
0      0  ...      0      0      0      0      0
1      0  ...      0      0      0      0      0
2      0  ...      0      0      0      0      0
3      0  ...      0      0      0      0      0
4      0  ...      0      0      0      0      0
```

```
seeds_772  seeds_773  seeds_774  seeds_775  seeds_776
0      0      0      0      0      0
1      0      0      0      0      0
2      0      0      0      0      0
3      0      0      0      0      0
4      0      0      0      0      0
```

[5 rows x 3195 columns]

```
Target:
yield_0  yield_1  yield_2  yield_3  yield_4  yield_5  yield_6  yield_7  \
0      0      0      0      0      0      0      0      0
1      0      0      0      0      0      0      0      0
2      0      0      0      0      0      0      0      0
3      0      0      0      0      0      0      0      0
4      0      0      0      0      0      0      0      0
```

```
yield_8  yield_9  ...  yield_767  yield_768  yield_769  yield_770  \
0      0      0  ...      0      0      0      0
1      0      0  ...      0      0      0      0
2      0      0  ...      0      0      0      0
3      0      0  ...      0      0      0      0
4      0      0  ...      0      0      0      0
```

```
yield_771  yield_772  yield_773  yield_774  yield_775  yield_776
0      0      0      0      0      0      0
1      0      0      0      0      0      0
2      0      0      0      0      0      0
3      0      0      0      0      0      0
4      0      0      0      0      0      0
```

[5 rows x 777 columns]

Task 3: Modeling and Evaluation

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

X = df.drop(columns=df.filter(like='yield').columns)
y = df.filter(like='yield')

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a Linear Regression model
model = LinearRegression()

# Train the model on the training data
model.fit(X_train, y_train)

# LinearRegression
LinearRegression()

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model using Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse}')

Mean Squared Error: 0.0013361585157398892

print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)

X_train shape: (621, 3195)
y_train shape: (621, 777)

from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error

# Extract the first target variable for SVR
y_train_single = y_train.iloc[:, 0]
y_test_single = y_test.iloc[:, 0]

# Create a Support Vector Regressor model
svr_model = SVR()

# Train the model on the training data
svr_model.fit(X_train, y_train_single)

# Make predictions on the test set
svr_y_pred_single = svr_model.predict(X_test)

# Evaluate the model using Mean Squared Error (MSE)
svr_mse_single = mean_squared_error(y_test_single, svr_y_pred_single)
print(f'Support Vector Regressor Mean Squared Error: {svr_mse_single}')

Support Vector Regressor Mean Squared Error: 0.01033312492949816

from sklearn.multioutput import MultiOutputRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

# Create a Decision Tree Regressor model
tree_model = DecisionTreeRegressor()

# Wrap the Decision Tree Regressor model in a MultiOutputRegressor
multioutput_tree = MultiOutputRegressor(tree_model)

# Train the model on the training data
multioutput_tree.fit(X_train, y_train)

# Make predictions on the test set
tree_y_pred = multioutput_tree.predict(X_test)

# Evaluate the model using Mean Squared Error (MSE)
tree_mse = mean_squared_error(y_test, tree_y_pred)
print(f'MultiOutput Decision Tree Regressor Mean Squared Error: {tree_mse}')

MultiOutput Decision Tree Regressor Mean Squared Error: 0.0012952512952512953

from sklearn.tree import DecisionTreeRegressor
from sklearn.multioutput import MultiOutputRegressor
from sklearn.metrics import mean_squared_error

# Create a Decision Tree Regressor model
dt_base_model = DecisionTreeRegressor(random_state=42)

# Wrap the base model in a MultiOutputRegressor
multioutput_dt_model = MultiOutputRegressor(dt_base_model)

# Train the model on the training data
multioutput_dt_model.fit(X_train, y_train)

# Make predictions on the test set
dt_y_pred = multioutput_dt_model.predict(X_test)

# Evaluate the model using Mean Squared Error (MSE)
dt_mse = mean_squared_error(y_test, dt_y_pred)
print(f'Decision Tree Mean Squared Error: {dt_mse}')

Decision Tree Mean Squared Error: 0.0012952512952512953

from sklearn.ensemble import RandomForestRegressor
from sklearn.multioutput import MultiOutputRegressor
from sklearn.metrics import mean_squared_error

# Create a Random Forest Regressor model

```

```

rf_base_model = RandomForestRegressor(random_state=42)

# Wrap the base model in a MultiOutputRegressor
multioutput_rf_model = MultiOutputRegressor(rf_base_model)

# Train the model on the training data
multioutput_rf_model.fit(X_train, y_train)

# Make predictions on the test set
rf_y_pred = multioutput_rf_model.predict(X_test)

# Evaluate the model using Mean Squared Error (MSE)
rf_mse = mean_squared_error(y_test, rf_y_pred)
print(f'Random Forest Mean Squared Error: {rf_mse}')

Random Forest Mean Squared Error: 0.0012888963138963142

# Import seaborn as sns
# Import matplotlib.pyplot as plt

# Assuming df is your DataFrame
# correlation_matrix = df.corr()

# Set up the matplotlib figure
# plt.figure(figsize=(12, 10))

# Create a heatmap using Seaborn
# sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=.5)

# Show the plot
# plt.show()

```

Session crashed when executed the above

```

File "<ipython-input-28-3cedb23d147e>", line 1
    Session crashed when executed the above
    ^
SyntaxError: invalid syntax

```

Task 4: Hyperparameter Tuning

```

from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Print the shapes to ensure consistency
print(f'X_train shape: {X_train.shape}, y_train shape: {y_train.shape}')
print(f'X_test shape: {X_test.shape}, y_test shape: {y_test.shape}')

X_train shape: (621, 3195), y_train shape: (621, 777)
X_test shape: (156, 3195), y_test shape: (156, 777)

# Last column as the target
y_train_single = y_train.iloc[:, -1]
y_test_single = y_test.iloc[:, -1]

from sklearn.model_selection import GridSearchCV

from sklearn.multioutput import MultiOutputRegressor
from sklearn.svm import SVR

# Create a Support Vector Regressor model
svr_model = SVR()

# Wrap the SVR model in MultiOutputRegressor
multioutput_svr = MultiOutputRegressor(svr_model)

# Define the hyperparameter grid to search
param_grid = {
    'estimator__C': [0.1, 1, 10],
    'estimator__kernel': ['linear', 'rbf', 'poly'],
    'estimator__gamma': ['scale', 'auto']
}

# Create GridSearchCV
grid_search = GridSearchCV(estimator=multioutput_svr, param_grid=param_grid, scoring='neg_mean_squared_error', cv=5)

# Fit the model to the training data
grid_search.fit(X_train, y_train)

# Get the best hyperparameters
best_params = grid_search.best_params_

# Print the best hyperparameters
print(f'Best Hyperparameters: {best_params}')

# Make predictions on the test set using the best model
svr_y_pred = grid_search.predict(X_test)

# Evaluate the model using Mean Squared Error (MSE)
svr_mse = mean_squared_error(y_test, svr_y_pred)
print(f'Support Vector Regressor Mean Squared Error: {svr_mse}')

Best Hyperparameters: {'estimator__C': 0.1, 'estimator__gamma': 'scale', 'estimator__kernel': 'linear'}
Support Vector Regressor Mean Squared Error: 0.003515131926795013

```

Task 5: Explainable AI

```
!pip install shap==0.40.0
```

```

Collecting shap==0.40.0
  Downloading shap-0.40.0.tar.gz (371 kB)
    371.7/371.7 kB 8.4 MB/s eta 0:00:00
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from shap==0.40.0) (1.25.2)

```

```
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from shap==0.40.0) (1.11.4)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from shap==0.40.0) (1.2.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from shap==0.40.0) (1.5.3)
Requirement already satisfied: tqdm>4.25.0 in /usr/local/lib/python3.10/dist-packages (from shap==0.40.0) (4.66.2)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/dist-packages (from shap==0.40.0) (23.2)
Collecting slicer==0.0.7 (from shap==0.40.0)
  Downloading slicer-0.0.7-py3-none-any.whl (14 kB)
Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-packages (from shap==0.40.0) (0.58.1)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-packages (from shap==0.40.0) (2.2.1)
Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba->shap==0.40.0) (0.41.1)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap==0.40.0) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap==0.40.0) (2023.4)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap==0.40.0) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap==0.40.0) (3.2.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas->shap==0.40.0) (1.16.0)
Building wheels for collected packages: shap
  Building wheel for shap (pyproject.toml) ... done
  Created wheel for shap: filename=shap-0.40.0-cp310-cp310-linux_x86_64.whl size=515069 sha256=4032d5df8578aec2c53d025e335c1ea159b17d4ae4ebf5ea11165a6535cc97ad
  Stored in directory: /root/.cache/pip/wheels/33/28/e3/62a9dc612c58c1b8d1c16fa51e64941bbb38ac8a6decbad39c
Successfully built shap
Installing collected packages: slicer, shap
Successfully installed shap-0.40.0 slicer-0.0.7

!pip install shap

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import shap
import numpy as np
import matplotlib.pyplot as plt

import pandas as pd

# Provide the path to your dataset file
file_path = "/content/WildDBlueberryPollinationSimulationData.csv"

# Read the dataset into a Pandas DataFrame
df = pd.read_csv(file_path)

# Check the column names
print(df.columns)

# Assuming 'yield' is one of the column names
# If 'yield' is present, define features and target variables
if 'yield' in df.columns:
    features = df.drop('yield', axis=1, errors='ignore') # Drop 'yield' column if present
    target = df['yield']

    # Create a Linear Regression model
    linear_model = LinearRegression()

    # Assuming you have your dataset loaded and split into features and target variables
    X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2, random_state=42)

    # Create and fit the linear regression model
    linear_model.fit(X_train, y_train)

    # Define a callable predict function
    def predict_fn(X):
        return linear_model.predict(X)

    # Create an explainer object using KernelExplainer
    explainer = shap.KernelExplainer(predict_fn, X_train)

    # Explain the predictions on a single instance (you can choose any index)
    sample_index = 0
    shap_values = explainer.shap_values(X_train.iloc[sample_index])

    # Reshape shap_values to have two dimensions
    shap_values = shap_values.reshape(1, -1)

    # Force plot for the chosen instance
    shap.force_plot(explainer.expected_value, shap_values, X_train.iloc[sample_index])

    # Summary plot for all instances
    shap.summary_plot(shap_values, X_train, plot_type='bar') # Use plot_type='bar' to show feature importance
else:
    print("The 'yield' column is not present in the DataFrame.")

Requirement already satisfied: shap in /usr/local/lib/python3.10/dist-packages (0.44.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from shap) (1.25.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from shap) (1.11.4)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from shap) (1.2.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from shap) (1.5.3)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.10/dist-packages (from shap) (4.66.2)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/dist-packages (from shap) (23.2)
Requirement already satisfied: slicer==0.0.7 in /usr/local/lib/python3.10/dist-packages (from shap) (0.0.7)
Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-packages (from shap) (0.58.1)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-packages (from shap) (2.2.1)
Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba->shap) (0.41.1)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2023.4)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (3.2.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas->shap) (1.16.0)
X does not have valid feature names, but LinearRegression was fitted with feature names
WARNING:shap:Using 621 background data samples could cause slower run times. Consider using shap.sample(data, K) or shap.kmeans(data, K) to summarize the background as K samples.
Index(['Row#', 'clonesize', 'honeybee', 'bumbles', 'andrena', 'osmia',
      'MaxOfUpperTRange', 'MinOfUpperTRange', 'AverageOfUpperTRange',
      'MaxOfLowerTRange', 'MinOfLowerTRange', 'AverageOfLowerTRange',
      'RainingDays', 'AverageRainingDays', 'fruitset', 'fruitmass', 'seeds',
      'yield'],
      dtype='object')
X does not have valid feature names, but LinearRegression was fitted with feature names
X does not have valid feature names, but LinearRegression was fitted with feature names
```

MaxOfUpperTRange

AverageOfLowerTRange

AverageOfLowerTRange